Chapter 3. Using programmable I/O (PIO)

3.1. What is Programmable I/O (PIO)?

Programmable I/O (PIO) is a new piece of hardware developed for RP2040. It allows you to create new types of (or additional) hardware interfaces on your RP2040-based device. If you've looked at fixed peripherals on a microcontroller, and thought "I want to add 4 more UARTs", or "I'd like to output DPI video", or even "I need to communicate with this cursed serial device I found on AliExpress, but no machine has hardware support", then you will have fun with this chapter.

PIO hardware is described extensively in chapter 3 of the RP2040 Datasheet. This is a companion to that text, focussing on how, when and why to use PIO in your software. To start, we're going to spend a while discussing why I/O is hard, what the current options are, and what PIO does differently, before diving into some software tutorials. We will also try to illuminate some of the more important parts of the hardware along the way, but will defer to the datasheet for full explanations.



You can skip to the first software tutorial if you'd prefer to dive straight in.

3.1.1. Background

Interfacing with other digital hardware components is hard. It often happens at very high frequencies (due to amounts of data that need to be transferred), and has very exact timing requirements.

3.1.2. I/O Using dedicated hardware on your PC

Traditionally, on your desktop or laptop computer, you have one option for hardware interfacing. Your computer has high speed USB ports, HDMI outputs, PCIe slots, SATA drive controllers etc. to take care of the tricky and time sensitive business of sending and receiving ones and zeros, and responding with minimal latency or interruption to the graphics card, hard drive etc. on the other end of the hardware interface.

The custom hardware components take care of specific tasks that the more general multi-tasking CPU is not designed for. The operating system drivers perform higher level management of what the hardware components do, and coordinate data transfers via DMA to/from memory from the controller and receive IRQs when high level tasks need attention. These interfaces are purpose-built, and if you have them, you should use them.

3.1.3. I/O Using dedicated hardware on your Raspberry Pi or microcontroller

Not so common on PCs: your Raspberry Pi or microcontroller is likely to have dedicated hardware on chip for managing UART, I2C, SPI, PWM, I2S, CAN bus and more over general purpose I/O pins (GPIOs). Like USB controllers (also found on some microcontrollers, including the RP2040 on Raspberry Pi Pico), I2C and SPI are general purpose buses which connect to a wide variety of external hardware, using the same piece of on-chip hardware. This includes sensors, external flash, EEPROM and SRAM memories, GPIO expanders, and more, all of them widely and cheaply available. Even HDMI uses I2C to communicate video timings between Source and Sink, and there is probably a microcontroller embedded in your TV to handle this.

These protocols are simpler to integrate into very low-cost devices (i.e. not the host), due to their relative simplicity and

modest speed. This is important for chips with mostly analogue or high-power circuitry: the silicon fabrication techniques used for these chips do not lend themselves to high speed or gate count, so if your switchmode power supply controller has some serial configuration interface, it is likely to be something like I2C. The number of traces routed on the circuit board, the number of pins required on the device package, and the PCB technology required to maintain signal integrity are also factors in the choice of these protocols. A microcontroller needs to communicate with these devices to be part of a larger *embedded system*.

This is all very well, but the area taken up by these individual serial peripherals, and the associated cost, often leaves you with a limited menu. You may end up paying for a bunch of stuff you don't need, and find yourself without enough of what you really want. Of course you are out of luck if your microcontroller does not have dedicated hardware for the type of hardware device you want to attach (although in some cases you may be able to bridge over USB, I2C or SPI at the cost of buying external hardware).

3.1.4. I/O Using software control of GPIOs ("bit-banging")

The third option on your Raspberry Pi or microcontroller — any system with GPIOs which the processor(s) can access easily — is to use the CPU to wiggle (and listen to) the GPIOs at dizzyingly high speeds, and hope to do so with sufficiently correct timing that the external hardware still understands the signals.

As a bit of background it is worth thinking about types of hardware that you might want to interface, and the approximate signalling speeds involved:

Table 4. Types of hardware

Interface Speed	Interface
1-10Hz	Push buttons, indicator LEDs
300Hz	HDMI CEC
10-100kHz	Temperature sensors (DHT11), one-wire serial
<100kHz	I2C Standard mode
22-100+kHz	PCM audio
300+kHz	PWM audio
400-1200kHz	WS2812 LED string
10-3000kHz	UART serial
12MHz	USB Full Speed
1-100MHz	SPI
20-300MHz	DPI/VGA video
480MHz	USB High Speed
10-4000MHz	Ethernet LAN
12-4000MHz	SD card
250-20000MHz	HDMI/DVI video

"Bit-Banging" (i.e. using the processor to hammer out the protocol via the GPIOs) is very hard. The processor isn't really designed for this. It has other work to do... for slower protocols you might be able to use an IRQ to wake up the processor from what it was doing fast enough (though latency here is a concern) to send the next bit(s). Indeed back in the early days of PC sound it was not uncommon to set a hardware timer interrupt at 11kHz and write out one 8-bit PCM sample every interrupt for some rather primitive sounding audio!

Doing that on a PC nowadays is laughed at, even though they are many order of magnitudes faster than they were back then. As processors have become faster in terms of overwhelming number-crunching brute force, the layers of software and hardware between the processor and the outside world have also grown in number and size. In response to the growing distance between processors and memory, PC-class processors keep many hundreds of instructions in-flight

on a single core at once, which has drawbacks when trying to switch rapidly between hard real time tasks. However, IRQ-based bitbanging can be an effective strategy on simpler embedded systems.

Above certain speeds — say a factor of 1000 below the processor clock speed — IRQs become impractical, in part due to the timing uncertainty of actually *entering* an interrupt handler. The alternative when "bit-banging" is to sit the processor in a carefully timed loop, often painstakingly written in assembly, trying to make sure the GPIO reading and writing happens on the exact cycle required. This is really really hard work if indeed possible at all. Many heroic hours and likely thousands of GitHub repositories are dedicated to the task of doing such things (a large proportion of them for LED strings).

Additionally of course, your processor is now busy doing the "bit-banging", and cannot be used for other tasks. If your processor is interrupted even for a few microseconds to attend to one of the hard peripherals it is also responsible for, this can be fatal to the timing of any bit-banged protocol. The greater the ratio between protocol speed and processor speed, the more cycles your processor will spend uselessly idling in between GPIO accesses. Whilst it is eminently possible to drive a 115200 baud UART output using only software, this has a cost of >10,000 cycles per byte if the processor is running at 133MHz, which may be poor investment of those cycles.

Whilst dealing with something like an LED string is possible using "bit-banging", once your hardware protocol gets faster to the point that it is of similar order of magnitude to your system clock speed, there is really not much you can hope to do. The main case where software GPIO access is the best choice is LEDs and push buttons.

Therefore you're back to custom hardware for the protocols you know up front you are going to want (or more accurately, the chip designer thinks you might need).

3.1.5. Programmable I/O Hardware using FPGAs and CPLDs

A field-programmable gate array (FPGA), or its smaller cousin, the complex programmable logic device (CPLD), is in many ways the perfect solution for tailor-made I/O requirements, whether that entails an unusual type or unusual mixture of interfaces. FPGAs are chips with a configurable logic fabric — effectively a sea of gates and flipflops, some other special digital function blocks, and a routing fabric to connect them — which offer the same level of design flexibility available to chip designers. This brings with it all the advantages of dedicated I/O hardware:

- Absolute precision of protocol timing (within limitations of your clock source)
- Capable of very high I/O throughput
- Offload simple, repetitive calculations that are part of the I/O standard (checksums)
- Present a simpler interface to host software; abstract away details of the protocol, and handle these details internally.

The main drawback of FPGAs in embedded systems is their cost. They also present a very unfamiliar programming model to those well-versed in embedded software: you are not programming at all, but rather designing digital hardware. One you have your FPGA you will still need some other processing element in your system to run control software, unless you are using an FPGA expensive enough to either fit a soft CPU core, or contain a hardened CPU core alongside the FPGA fabric.

eFPGAs (embedded FPGAs) are available in some microcontrollers: a slice of FPGA logic fabric integrated into a more conventional microcontroller, usually with access to some GPIOs, and accessible over the system bus. These are attractive from a system integration point of view, but have a significant area overhead compared with the usual serial peripherals found on a microcontroller, so either increase the cost and power dissipation, or are very limited in size. The issue of programming complexity still remains in eFPGA-equipped systems.

3.1.6. Programmable I/O Hardware using PIO

The PIO subsystem on RP2040 allows you to write small, simple programs for what are called *PIO state machines*, of which RP2040 has eight split across two PIO *instances*. A state machine is responsible for setting and reading one or more GPIOs, buffering data to or from the processor (or RP2040's ultra-fast DMA subsystem), and notifying the processor, via IRQ or polling, when data or attention is needed.

These programs operate with cycle accuracy at up to system clock speed (or the program clocks can be divided down to run at slower speeds for less frisky protocols).

PIO state machines are much more compact than the general-purpose Cortex-M0+ processors on RP2040. In fact, they are similar in size (and therefore cost) to a standard SPI peripheral, such as the PL022 SPI also found on RP2040, because much of their area is spent on components which are common to all serial peripherals, like FIFOs, shift registers and clock dividers. The instruction set is small and regular, so not much silicon is spent on decoding the instructions. There is no need to feel guilty about dedicating a state machine solely to a single I/O task, since you have 8 of them!

In spite of this, a PIO state machine gets a lot more done in one cycle than a Cortex-M0+ when it comes to I/O: for example, sampling a GPIO value, toggling a clock signal and pushing to a FIFO all in one cycle, every cycle. The trade-off is that a PIO state machine is not remotely capable of running general purpose software. As we shall see though, programming a PIO state machine is quite familiar for anyone who has written assembly code before, and the small instruction set should be fairly quick to pick up for those who haven't.

For simple hardware protocols - such as PWM or duplex SPI - a single PIO state machine can handle the task of implementing the hardware interface all on its own. For more involved protocols such as SDIO or DPI video you may end up using two or three.



TIP

If you are ever tempted to "bit-bang" a protocol on RP2040, don't! Use the PIO instead. Frankly this is true for anything that repeatedly reads or writes from GPIOs, but certainly anything which aims to transfer data.

3.2. Getting started with PIO

It is possible to write PIO programs both within the C++ SDK and directly from MicroPython.

Additionally the future intent is to add APIs to trivially have new UARTs, PWM channels etc created for you, using a menu of pre-written PIO programs, but for now you'll have to follow along with example code and do that yourself.

3.2.1. A First PIO Application

Before getting into all of the fine details of the PIO assembly language, we should take the time to look at a small but complete application which:

- 1. Loads a program into a PIO's instruction memory
- 2. Sets up a PIO state machine to run the program
- 3. Interacts with the state machine once it is running.

The main ingredients in this recipe are:

- A PIO program
- · Some software, written in C, to run the whole show
- A CMake file describing how these two are combined into a program image to load onto a RP2040-based development board

TIP

The code listings in this section are all part of a complete application on GitHub, which you can build and run. Just click the link above each listing to go to the source. In this section we are looking at the pio/hello_pio example in pico-examples. You might choose to build this application and run it, to see what it does, before reading through this section.

NOTE

The focus here is on the main moving parts required to use a PIO program, not so much on the PIO program itself. This is a lot to take in, so we will stay high-level in this example, and dig in deeper on the next one.

3.2.1.1. PIO Program

This is our first PIO program listing. It's written in PIO assembly language.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.pio Lines 7 - 15

```
7 .program hello
9 ; Repeatedly get one word of data from the TX FIFO, stalling when the FIFO is
10 ; empty. Write the least significant bit to the OUT pin group.
11
12 loop:
    pull
13
     out pins, 1
14
15
   jmp loop
```

The pull instruction takes one data item from the transmit FIFO buffer, and places it in the output shift register (OSR). Data moves from the FIFO to the OSR one word (32 bits) at a time. The OSR is able to shift this data out, one or more bits at a time, to further destinations, using an out instruction.

FIFOs?

FIFOs are data queues, implemented in hardware. Each state machine has two FIFOs, between the state machine and the system bus, for data travelling out of (TX) and into (RX) the chip. Their name (first in, first out) comes from the fact that data appears at the FIFO's output in the same order as it was presented to the FIFO's input.

The out instruction here takes one bit from the data we just pull-ed from the FIFO, and writes that data to some pins. We will see later how to decide which pins these are.

The jmp instruction jumps back to the loop: label, so that the program repeats indefinitely. So, to sum up the function of this program: repeatedly take one data item from a FIFO, take one bit from this data item, and write it to a pin.

Our .pio file also contains a helper function to set up a PIO state machine for correct execution of this program:

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.pio Lines 18 - 33

```
18 static inline void hello_program_init(PIO pio, uint sm, uint offset, uint pin) {
     pio_sm_config c = hello_program_get_default_config(offset);
21
      // Map the state machine's OUT pin group to one pin, namely the `pin`
22
      // parameter to this function.
       sm_config_set_out_pins(&c, pin, 1);
```

```
// Set this pin's GPIO function (connect PIO to the pad)
pio_gpio_init(pio, pin);

// Set the pin direction to output at the PIO
pio_sm_set_consecutive_pindirs(pio, sm, pin, 1, true);

// Load our configuration, and jump to the start of the program
pio_sm_init(pio, sm, offset, &c);

// Set the state machine running
pio_sm_set_enabled(pio, sm, true);

// Set the state machine running
```

Here the main thing to set up is the GPIO we intend to output our data to. There are three things to consider here:

- 1. The state machine needs to be told which GPIO or GPIOs to output to. There are four different pin groups which are used by different instructions in different situations; here we are using the out pin group, because we are just using an out instruction.
- 2. The GPIO also needs to be told that PIO is in control of it (GPIO function select)
- 3. If we are using the pin for output only, we need to make sure that PIO is driving the *output enable* line high. PIO can drive this line up and down programmatically using e.g. an out pindirs instruction, but here we are setting it up before starting the program.

3.2.1.2. C Program

PIO won't do anything until it's been configured properly, so we need some software to do that. The PIO file we just looked at — hello.pio — is converted automatically (we will see later how) into a header containing our assembled PIO program binary, any helper functions we included in the file, and some useful information about the program. We include this as hello.pio.h.

 ${\it Pico~Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/hello.com/raspberrypi/pico-examples/blob/master/pico-examples/blob/master/pico-examples/blob/master/pico-examples/blob/master/pico-examples/blob/master/pico-examples/blob/master/pico-examples/blob/ma$

```
2 * Copyright (c) 2020 Raspberry Pi (Trading) Ltd.
3
4 * SPDX-License-Identifier: BSD-3-Clause
5 */
7 #include "pico/stdlib.h"
8 #include "hardware/pio.h"
9 // Our assembled program:
10 #include "hello.pio.h"
11
12 int main() {
13 #ifndef PICO_DEFAULT_LED_PIN
14 #warning pio/hello_pio example requires a board with a regular LED
15 #else
16
   // Choose which PIO instance to use (there are two instances)
17
     PIO pio = pio0;
18
    // Our assembled program needs to be loaded into this PIO's instruction
19
20
      // memory. This SDK function will find a location (offset) in the
21
      // instruction memory where there is enough space for our program. We need
      // to remember this location!
      uint offset = pio_add_program(pio, &hello_program);
24
25
      // Find a free state machine on our chosen PIO (erroring if there are
      // none). Configure it to run our program, and start it, using the
27
       // helper function we included in our .pio file.
28
       uint sm = pio_claim_unused_sm(pio, true);
       hello_program_init(pio, sm, offset, PICO_DEFAULT_LED_PIN);
29
```

```
30
      // The state machine is now running. Any value we push to its TX FIFO will
31
32
      // appear on the LED pin.
33
      while (true) {
          // Blink
34
35
           pio_sm_put_blocking(pio, sm, 1);
36
          sleep_ms(500);
37
           // Blonk
           pio_sm_put_blocking(pio, sm, 0);
38
39
           sleep_ms(500);
40
41 #endif
42 }
```

You might recall that RP2040 has two PIO blocks, each of them with four state machines. Each PIO block has a 32-slot instruction memory which is visible to the four state machines in the block. We need to load our program into this instruction memory before any of our state machines can run the program. The function pio_add_program() finds free space for our program in a given PIO's instruction memory, and loads it.

32 Instructions?

This may not sound like a lot, but the PIO instruction set can be *very* dense once you fully explore its features. A perfectly serviceable UART transmit program can be implemented in four instructions, as shown in the pio/uart_tx example in pico-examples. There are also a couple of ways for a state machine to execute instructions from other sources — like directly from the FIFOs — which you can read all about in the RP2040 Datasheet.

Once the program is loaded, we find a free state machine and tell it to run our program. There is nothing stopping us from ordering multiple state machines to run the same program. Likewise, we could instruct each state machine to run a *different* program, provided they all fit into the instruction memory at once.

We're configuring this state machine to output its data to the LED on your Raspberry Pi Pico board. If you have already built and run the program, you probably noticed this already!

At this point, the state machine is running autonomously. The state machine will immediately *stall*, because it is waiting for data in the TX FIFO, and we haven't provided any. The processor can push data directly into the state machine's TX FIFO using the pio_sm_put_blocking() function. (_blocking because this function stalls the processor when the TX FIFO is full.) Writing a 1 will turn the LED on, and writing a 0 will turn the LED off.

3.2.1.3. CMake File

We have two lovely text files sat on our computer, with names ending with .pio and .c, but they aren't doing us much good there. A CMake file describes how these are built into a binary suitable for loading onto your Raspberry Pi Pico or other RP2040-based board.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/hello_pio/CMakeLists.txt

```
11
12 pico_add_extra_outputs(hello_pio)
13
14 # add url via pico_set_program_url
15 example_auto_set_url(hello_pio)
```

- add_executable(): Declare that we are building a program called hello_pio
- pico_generate_pio_header(): Declare that we have a PIO program, hello.pio, which we want to be built into a C header for use with our program
- target_sources(): List the source code files for our hello_pio program. In this case, just one C file.
- target_link_libraries(): Make sure that our program is built with the PIO hardware API, so we can call functions like pio_add_program() in our C file.
- pico_add_extra_outputs(): By default we just get an .elf file as the build output of our app. Here we declare we also
 want extra build formats, like a .uf2 file which can be dragged and dropped directly onto a Raspberry Pi Pico
 attached over USB.

Assuming you already have pico-examples and the SDK installed on your machine, you can run

```
$ mkdir build
$ cd build
$ cmake ..
$ make hello_pio
```

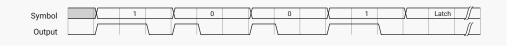
To build this program.

3.2.2. A Real Example: WS2812 LEDs

The WS2812 LED (sometimes sold as NeoPixel) is an addressable RGB LED. In other words, it's an LED where the red, green and blue components of the light can be individually controlled, and it can be connected in such a way that many WS2812 LEDs can be controlled individually, with only a single control input. Each LED has a pair of power supply terminals, a serial data input, and a serial data output.

When serial data is presented at the LED's input, it takes the first three bytes for itself (red, green, blue) and the remainder is passed along to its serial data output. Often these LEDs are connected in a single long chain, each LED connected to a common power supply, and each LED's data output connected through to the next LED's input. A long burst of serial data to the first in the chain (the one with its data input unconnected) will deposit three bytes of RGB data in each LED, so their colour and brightness can be individually programmed.

Figure 3. WS2812 line format. Wide positive pulse for 1, narrow positive pulse for 0, very long negative pulse for latch enable



Unfortunately the LEDs receive and retransmit serial data in quite an unusual format. Each bit is transferred as a positive pulse, and the width of the pulse determines whether it is a 1 or a 0 bit. There is a family of WS2812-like LEDs available, which often have slightly different timings, and demand precision. It is possible to bit-bang this protocol, or to write canned bit patterns into some generic serial peripheral like SPI or I2S to get firmer guarantees on the timing, but there is still some software complexity and cost associated with generating the bit patterns.

Ideally we would like to have all of our CPU cycles available to generate colour patterns to put on the lights, or to handle any other responsibilities the processor may have in the *embedded system* the LEDs are connected to.



TIP

Once more, this section is going to discuss a real, complete program, that you can build and run on your Raspberry Pi Pico. Follow the links above the program listings if you'd prefer to build the program yourself and run it, before going through it in detail. This section explores the pio/ws2812 example in pico-examples.

3.2.2.1. PIO Program

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/ws2812.pio Lines 7 - 26

```
7 .program ws2812
8 .side_set 1
10 .define public T1 2
11 .define public T2 5
12 .define public T3 3
13
14 .lang_opt python sideset_init = pico.PIO.OUT_HIGH
15 .lang_opt python out_init = pico.PIO.OUT_HIGH
16 .lang_opt python out_shiftdir = 1
17
18 .wrap_target
19 bitloop:
20 out x, 1
                   side 0 [T3 - 1]; Side-set still takes place when instruction stalls
21 jmp !x do_zero side 1 [T1 - 1]; Branch on the bit we shifted out. Positive pulse
22 do_one:
23 jmp bitloop side 1 [T2 - 1]; Continue driving high, for a long pulse
24 do_zero:
25 nop
                   side 0 [T2 - 1] ; Or drive low, for a short pulse
26 .wrap
```

The previous example was a bit of a whistle-stop tour of the anatomy of a PIO-based application. This time we will dissect the code line-by-line. The first line tells the assembler that we are defining a program named ws2812:

```
.program ws2812
```

We can have multiple programs in one .pio file (and you will see this if you click the GitHub link above the main program listing), and each of these will have its own .program directive with a different name. The assembler will go through each program in turn, and all the assembled programs will appear in the output file.

Each PIO instruction is 16 bits in size. Generally, 5 of those bits in each instruction are used for the "delay" which is usually 0 to 31 cycles (after the instruction completes and before moving to the next instruction). If you have read the PIO chapter of the RP2040 Datasheet, you may have already know that these 5 bits can be used for a different purpose:

```
.side set 1
```

This directive .side_set 1 says we're stealing one of those delay bits to use for "side-set". The state machine will use this bit to drive the values of some pins, once per instruction, in addition to what the instructions are themselves doing. This is very useful for high frequency use cases (e.g. pixel clocks for DPI panels), but also for shrinking program size, to fit into the shared instruction memory.

Note that stealing one bit has left our delay range from 0-15 (4 bits), but that is quite natural because you rarely want to mix side-set with lower frequency stuff. Because we didn't say .side_set 1 opt, which would mean the side-set is optional (at the cost of another bit to say *whether* the instruction does a side-set), we have to specify a side-set value for *every* instruction in the program. This is the side N you will see on each instruction in the listing.

```
.define public T1 2
.define public T2 5
.define public T3 3
```

.define lets you declare constants. The public keyword means that the assembler will also write out the value of the define in the output file for use by other software: in the context of the SDK, this is a #define. We are going to use T1, T2 and T3 in calculating the delay cycles on each instruction.

```
.lang_opt python
```

This is used to specify some PIO hardware defaults as used by the MicroPython PIO library. We don't need to worry about them in the context of SDK applications.

```
.wrap_target
```

We'll ignore this for now, and come back to it later, when we meet its friend .wrap.

```
bitloop:
```

This is a label. A label tells the assembler that this point in your code is interesting to you, and you want to refer to it later by name. Labels are mainly used with jmp instructions.

```
out x, 1 side 0 [T3 - 1] ; Side-set still takes place when instruction stalls
```

Finally we reach a line with a PIO instruction. There is a lot to see here.

- This is an out instruction. out takes some bits from the *output shift register* (OSR), and writes them somewhere else. In this case, the OSR will contain pixel data destined for our LEDs.
- [T3 1] is the number of delay cycles (T3 minus 1). T3 is a constant we defined earlier.
- x (one of two scratch registers; the other imaginatively called y) is the destination of the write data. State machines use their scratch registers to hold and compare temporary data.
- side 0: Drive low (0) the pin configured for side-set.
- Everything after the ; character is a comment. Comments are ignored by the assembler: they are just notes for humans to read.

Output Shift Register

The OSR is a staging area for data entering the state machine through the TX FIFO. Data is pulled from the TX FIFO into the OSR one 32-bit chunk at a time. When an out instruction is executed, the OSR can break this data into smaller pieces by *shifting* to the left or right, and sending the bits that drop off the end to one of a handful of different destinations, such as the pins.

The amount of data to be shifted is encoded by the out instruction, and the *direction* of the shift (left or right) is configured ahead of time. For full details and diagrams, see the RP2040 Datasheet.

So, the state machine will do the following operations when it executes this instruction:

- 1. Set 0 on the side-set pin (this happens even if the instruction stalls because no data is available in the OSR)
- 2. Shift one bit out of the OSR into the x register. The value of the x register will be either 0 or 1.
- 3. Wait T3 1 cycles after the instruction (I.e. the whole thing takes T3 cycles since the instruction itself took a cycle). Note that when we say cycle, we mean state machine execution cycles: a state machine can be made to execute at a slower rate than the system clock, by configuring its *clock divider*.

Let's look at the next instruction in the program.

```
jmp !x do_zero side 1 [T1 - 1] ; Branch on the bit we shifted out. Positive pulse
```

- 1. side 1 on the side-set pin (this is the leading edge of our pulse)
- 2. If x == 0 then go to the instruction labelled do_zero, otherwise continue on sequentially to the next instruction
- 3. We delay T1 1 after the instruction (whether the branch is taken or not)

Let's look at what our output pin has done so far in the program.



Figure 4. The state machine drives the line low for time T1 as it shifts out one data bit from the OSR, and then high for time T2 whilst branching on the value of the bit.

The pin has been low for time T3, and high for time T1. If the x register is 1 (remember this contains our 1 bit of pixel data) then we will fall through to the instruction labelled do_one:

```
do_one:

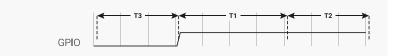
jmp bitloop side 1 [T2 - 1] ; Continue driving high, for a long pulse
```

On this side of the branch we do the following:

- 1. side 1 on the side-set pin (continue the pulse)
- 2. jmp unconditionally back to bitloop (the label we defined earlier, at the top of the program); the state machine is done with this data bit, and will get another from its OSR
- 3. Delay for T2 1 cycles after the instruction

The waveform at our output pin now looks like this:

Figure 5. On a one data bit, the line is driven low for time T3, high for time T1, then high for an additional time T2



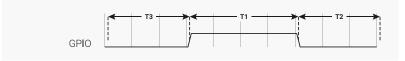
This accounts for the case where we shifted a 1 data bit into the x register. For a 0 bit, we will have jumped over the last instruction we looked at, to the instruction labelled do_zero:

```
do_zero:
nop side 0 [T2 - 1] ; Or drive low, for a short pulse
```

- 1. side 0 on the side-set pin (the trailing edge of our pulse)
- 2. nop means no operation. We don't have anything else we particularly want to do, so waste a cycle
- 3. The instruction takes T2 cycles in total

For the x == 0 case, we get this on our output pin:

Figure 6. On a zero data bit, the line is driven low for time T3, high for time T1, then low again for time T1



The final line of our program is this:

.wrap

This matches with the .wrap_target directive at the top of the program. Wrapping is a hardware feature of the state machine which behaves like a wormhole: you go in through the .wrap statement and appear at the .wrap_target zero cycles later, unless the .wrap is preceded immediately by a jmp whose condition is true. This is important for getting precise timing with programs that must run quickly, and often also saves you a slot in the instruction memory.



TIP

Often an explicit .wrap_target/.wrap pair is not necessary, because the default configuration produced by pioasm has an implicit wrap from the end of the program back to the beginning, if you didn't specify one.

NOPs

NOP, or no operation, means precisely that: do nothing! You may notice there is no nop instruction defined in the instruction set reference: nop is really a synonym for mov y, y in PIO assembly.

Why did we insert a nop in this example when we could have jmp-ed? Good question! It's a dramatic device we contrived so we could discuss nop and .wrap. Writing documentation is hard. In general, though, nop is useful when you need to perform a side-set and have nothing else to do, or you need a very slightly longer delay than is available on a single instruction.

It is hopefully becoming clear why our timings T1, T2, T3 are numbered this way, because what the LED string sees really is one of these two cases:

Figure 7. The line is initially low in the idle (latch) state, and the LED is waiting for the first rising edge. It sees our pulse timps in the order T1-T2-T3, until the very last T3, where it sees a much longer negative period once the state machine runs out of data.



This should look familiar if you refer back to Figure 3.

After thoroughly dissecting our program, and hopefully being satisfied that it will repeatedly send one well-formed data bit to a string of WS2812 LEDs, we're left with a question: where is the data coming from? This is more thoroughly explained in the RP2040 Datasheet, but the data that we are shifting out from the OSR came from the state machine's TX FIFO. The TX FIFO is a data buffer between the state machine and the rest of RP2040, filled either via direct poking from the CPU, or by the system DMA, which is much faster.

The out instruction shifts data out from the OSR, and zeroes are shifted in from the other end to fill the vacuum. Because the OSR is 32 bits wide, you will start getting zeroes once you have shifted out a total of 32 bits. There is a pull instruction which explicitly takes data from the TX FIFO and put it in the OSR (stalling the state machine if the FIFO is empty).

However, in the majority of cases it is simpler to configure *autopull*, a mode where the state machine automatically refills the OSR from the TX FIFO (an automatic pull) when a configured number of bits have been shifted out. Autopull happens in the background, in parallel with whatever else the state machine may be up to (in other words it has a cost of zero cycles). We'll see how this is configured in the next section.

3.2.2.2. State Machine Configuration

When we run pioasm on the .pio file we have been looking at, and ask it to spit out SDK code (which is the default), it will create some static variables describing the program, and a method ws2812_default_program_config which configures a PIO state machine based on user parameters, and the directives in the actual PIO program (namely the .side_set and .wrap in this case).

Of course how you configure the PIO SM when using the program is very much related to the program you have written. Rather than try to store a data representation off all that information, and parse it at runtime, for the use cases where you'd like to encapsulate setup or other API functions with your PIO program, you can embed code within the .pio file.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/ws2812.pio Lines 31 - 47

```
31 static inline void ws2812_program_init(PIO pio, uint sm, uint offset, uint pin, float freq,
  bool rgbw) {
32
33
       pio_gpio_init(pio, pin);
34
       pio_sm_set_consecutive_pindirs(pio, sm, pin, 1, true);
35
36
       pio_sm_config c = ws2812_program_get_default_config(offset);
37
       sm_config_set_sideset_pins(&c, pin);
       sm\_config\_set\_out\_shift(\&c, \ false, \ true, \ rgbw \ ? \ 32 \ : \ 24);
38
       sm\_config\_set\_fifo\_join(\&c, PIO\_FIFO\_JOIN\_TX);\\
39
40
41
       int cycles_per_bit = ws2812_T1 + ws2812_T2 + ws2812_T3;
42
       float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit);
43
       sm_config_set_clkdiv(&c, div);
44
45
       pio_sm_init(pio, sm, offset, &c);
46
       pio_sm_set_enabled(pio, sm, true);
47 }
```

In this case we are passing through code for the SDK, as requested by this line you will see if you click the link on the above listing to see the context:

```
% c-sdk {
```

We have here a function ws2812_program_init which is provided to help the user to instantiate an instance of the LED driver program, based on a handful of parameters:

pio

Which of RP2040's two PIO instances we are dealing with

SM

Which state machine on that PIO we want to configure to run the WS2812 program

offset

Where the PIO program was loaded in PIO's 5-bit program address space

pin

which GPIO pin our WS2812 LED chain is connected to

freq

The frequency (or rather baud rate) we want to output data at.

rgbw

True if we are using 4-colour LEDs (red, green, blue, white) rather than the usual 3.

Such that:

- pio_gpio_init(pio, pin); Configure a GPIO for use by PIO. (Set the GPIO function select.)
- pio_set_consecutive_pindirs(pio, sm, pin, 1, true); Sets the PIO pin direction of 1 pin starting at pin number pin to out
- pio_sm_config c = ws2812_program_default_config(offset); Get the default configuration using the generated function for this program (this includes things like the .wrap and .side_set configurations from the program). We'll modify this configuration before loading it into the state machine.
- sm_config_sideset_pins(&c, pin); Sets the side-set to write to pins starting at pin pin (we say starting at because if you had .side_set 3, then it would be outputting values on numbers pin, pin+1, pin+2)
- sm_config_out_shift(&c, false, true, rgbw ? 32 : 24); False for shift_to_right (i.e. we want to shift out MSB first). True for autopull. 32 or 24 for the number of bits for the autopull threshold, i.e. the point at which the state machine triggers a refill of the OSR, depending on whether the LEDs are RGB or RGBW.
- int cycles_per_bit = ws2812_T1 + ws2812_T2 + ws2812_T3; This is the total number of execution cycles to output a single bit. Here we see the benefit of .define public; we can use the T1 T3 values in our code.
- float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit); sm_config_clkdiv(&c, div); Slow the state machine's execution down, based on the system clock speed and the number of execution cycles required per WS2812 data bit, so that we achieve the correct bit rate.
- pio_sm_init(pio, sm, offset, &c); Load our configuration into the state machine, and go to the start address (
 offset)
- pio_sm_enable(pio, sm, true); And make it go now!

At this point the program will be stuck on the first out waiting for data. This is because we have autopull enabled, the OSR is initially empty, and there is no data to be pulled. The state machine refuses to continue until the first piece of data arrives in the FIFO.

As an aside, this last point sheds some light on the slightly cryptic comment at the start of the PIO program:

```
out x, 1 side \theta [T3 - 1]; Side-set still takes place when instruction stalls
```

This comment is giving us an important piece of context. We stall on this instruction initially, before the first data is added, and also every time we finish sending the last piece of data at the end of a long serial burst. When a state machine stalls, it does not continue to the next instruction, rather it will reattempt the current instruction on the next divided clock cycle. However, side-set still takes place. This works in our favour here, because we consequently always return the line to the idle (low) state when we stall.

3.2.2.3. C Program

The companion to the .pio file we've looked at is a .c file which drives some interesting colour patterns out onto a string of LEDs. We'll just look at the parts that are directly relevant to PIO.

Pico Examples: https://github.com/raspberryni/pico-examples/blob/master/pio/ws2812/ws2812.c Lines 25 - 27

```
25 static inline void put_pixel(uint32_t pixel_grb) {
26    pio_sm_put_blocking(pio0, 0, pixel_grb << 8u);
27 }
```

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/ws2812.c Lines 29 - 34

Here we are writing 32-bit values into the FIFO, one at a time, directly from the CPU. pio_sm_put_blocking is a helper method that waits until there is room in the FIFO before pushing your data.

You'll notice the << 8 in put_pixel(): remember we are shifting out starting with the MSB, so we want the 24-bit colour values at the top. This works fine for WGBR too, just that the W is always 0.

This program has a handful of colour patterns, which call our put_pixel helper above to output a sequence of pixel values:

```
50 void pattern_random(uint len, uint t) {
51    if (t % 8)
52        return;
53    for (int i = 0; i < len; ++i)
54        put_pixel(rand());
55 }</pre>
```

The main function loads the program onto a PIO, configures a state machine for 800 kbaud WS2812 transmission, and then starts cycling through the colour patterns randomly.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/ws2812.c Lines 84 - 108

```
91
        int sm = 0:
92
        uint offset = pio_add_program(pio, &ws2812_program);
93
94
        ws2812_program_init(pio, sm, offset, WS2812_PIN, 800000, IS_RGBW);
95
96
        int t = 0:
 97
        while (1) {
 98
           int pat = rand() % count_of(pattern_table);
            int dir = (rand() >> 30) & 1 ? 1 : -1;
99
            puts(pattern_table[pat].name);
100
            puts(dir == 1 ? "(forward)" : "(backward)");
101
            for (int i = 0; i < 1000; ++i) {
102
103
                pattern_table[pat].pat(NUM_PIXELS, t);
194
               sleep_ms(10);
105
                t += dir:
106
            }
107
        }
108 }
```

3.2.3. PIO and DMA (A Logic Analyser)

So far we have looked at writing data to PIO directly from the processor. This often leads to the processor spinning its wheels waiting for room in a FIFO to make a data transfer, which is not a good investment of its time. It also limits the total data throughput you can achieve.

RP2040 is equipped with a powerful *direct memory access* unit (DMA), which can transfer data for you in the background. Suitably programmed, the DMA can make quite long sequences of transfers without supervision. Up to one word per system clock can be transferred to or from a PIO state machine, which is, to be quite technically precise, more bandwidth than you can shake a stick at. The bandwidth is shared across all state machines, but you can use the full amount on *one* state machine.

Let's take a look at the logic_analyser example, which uses PIO to sample some of RP2040's own pins, and capture a logic trace of what is going on there, at full system speed.

 $Pico\ Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/logic_analyser/logic_analyser.c\ Lines\ 40-63$

```
40 void logic_analyser_init(PIO pio, uint sm, uint pin_base, uint pin_count, float div) {
41 // Load a program to capture n pins. This is just a single `in pins, n`
      // instruction with a wrap.
42
43
   uint16_t capture_prog_instr = pio_encode_in(pio_pins, pin_count);
44
    struct pio_program capture_prog = {
45
              .instructions = &capture_prog_instr,
              .length = 1,
46
47
              .origin = -1
48
      }:
49
      uint offset = pio_add_program(pio, &capture_prog);
      // Configure state machine to loop over this `in` instruction forever,
52
       // with autopush enabled.
53
       pio_sm_config c = pio_get_default_sm_config();
54
       sm_config_set_in_pins(&c, pin_base);
55
       sm_config_set_wrap(&c, offset, offset);
      sm_config_set_clkdiv(&c, div);
56
      // Note that we may push at a < 32 bit threshold if pin_count does not
57
      // divide 32. We are using shift-to-right, so the sample data ends up
58
59
      // left-justified in the FIFO in this case, with some zeroes at the LSBs.
      sm_config_set_in_shift(&c, true, true, bits_packed_per_word(pin_count));
       sm_config_set_fifo_join(&c, PIO_FIFO_JOIN_RX);
62
       pio_sm_init(pio, sm, offset, &c);
63 }
```

Our program consists only of a single in pins, <pin_count> instruction, with program wrapping and autopull enabled. Because the amount of data to be shifted is only known at runtime, and because the program is so short, we are generating the program dynamically here (using the pio_encode_ functions) instead of pushing it through pioasm. The program is wrapped in a data structure stating how big the program is, and where it must be loaded — in this case origin = -1 meaning "don't care".

Input Shift Register

The *input shift register* (ISR) is the mirror image of the OSR. Generally data flows through a state machine in one of two directions: System \rightarrow TX FIFO \rightarrow OSR \rightarrow Pins, or Pins \rightarrow ISR \rightarrow RX FIFO \rightarrow System. An in instruction shifts data into the ISR.

If you don't need the ISR's shifting ability — for example, if your program is output-only — you can use the ISR as a third scratch register. It's 32 bits in size, the same as X, Y and the OSR. The full details are in the RP2040 Datasheet.

We load the program into the chosen PIO, and then configure the input pin mapping on the chosen state machine so that its in pins instruction will see the pins we care about. For an in instruction we only need to worry about configuring the base pin, i.e. the pin which is the least significant bit of the in instruction's sample. The number of pins to be sampled is determined by the bit count parameter of the in pins instruction—it will sample *n* pins starting at the base we specified, and shift them into the ISR.

Pin Groups (Mapping)

We mentioned earlier that there are four pin groups to configure, to connect a state machine's internal data buses to the GPIOs it manipulates. A state machine accesses all pins within a group at once, and pin groups can overlap. So far we have seen the *out*, *side-set* and *in* pin groups. The fourth is *set*.

The out group is the pins affected by shifting out data from the OSR, using out pins or out pindirs, up to 32 bits at a time. The set group is used with set pins and set pindirs instructions, up to 5 bits at a time, with data that is encoded directly in the instruction. It's useful for toggling control signals. The side-set group is similar to the set group, but runs simultaneously with another instruction. Note: mov pin uses the in or out group, depending on direction.

Configuring the clock divider optionally slows down the state machine's execution: a clock divisor of n means 1 instruction will be executed per n system clock cycles. The default system clock frequency for SDK is 125MHz.

sm_config_set_in_shift sets the shift direction to rightward, enables autopush, and sets the autopush threshold to 32. The state machine keeps an eye on the total amount of data shifted into the ISR, and on the in which reaches or breaches a total shift count of 32 (or whatever number you have configured), the ISR contents, along with the new data from the in. goes straight to the RX FIFO. The ISR is cleared to zero in the same operation.

sm_config_set_fifo_join is used to manipulate the FIFOs so that the DMA can get more throughput. If we want to sample every pin on every clock cycle, that's a lot of bandwidth! We've finished describing how the state machine should be configured, so we use pio_sm_init to load the configuration into the state machine, and get the state machine into a clean initial state.

FIFO Joining

Each state machine is equipped with a FIFO going in each direction: the TX FIFO buffers data on its way out of the system, and the RX FIFO does the same for data coming in. Each FIFO has four data slots, each holding 32 bits of data. Generally you want FIFOs to be as deep as possible, so there is more slack time between the timing-critical operation of a peripheral, and data transfers from system agents which may be quite busy or have high access latency. However this comes with significant hardware cost.

If you are only using one of the two FIFOs — TX or RX — a state machine can pool its resources to provide a single FIFO with double the depth. The RP2040 Datasheet goes into much more detail, including how this mechanism actually works under the hood.

Our state machine is ready to sample some pins. Let's take a look at how we hook up the DMA to our state machine, and tell the state machine to start sampling once it sees some trigger condition.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/logic_analyser/logic_analyser.c Lines 65 - 87

```
65 void logic_analyser_arm(PIO pio, uint sm, uint dma_chan, uint32_t *capture_buf, size_t
  capture_size_words,
66
                           uint trigger_pin, bool trigger_level) {
67
      pio sm set enabled(pio. sm. false):
68
      // Need to clear _input shift counter_, as well as FIFO, because there may be
69
      // partial ISR contents left over from a previous run. sm_restart does this.
70
      pio_sm_clear_fifos(pio, sm);
71
      pio_sm_restart(pio, sm);
72
73
       dma_channel_config c = dma_channel_get_default_config(dma_chan);
74
       channel_config_set_read_increment(&c, false);
75
       channel_config_set_write_increment(&c, true);
76
       channel_config_set_dreq(&c, pio_get_dreq(pio, sm, false));
77
78
       dma_channel_configure(dma_chan, &c,
                           // Destination pointer
          capture_buf,
80
           &pio->rxf[sm],
                              // Source pointer
81
           capture_size_words, // Number of transfers
82
                               // Start immediately
83
       );
84
85
       pio_sm_exec(pio, sm, pio_encode_wait_gpio(trigger_level, trigger_pin));
86
       pio_sm_set_enabled(pio, sm, true);
87 }
```

We want the DMA to read from the RX FIFO on our PIO state machine, so every DMA read is from the same address. The *write* address, on the other hand, should increment after every DMA transfer so that the DMA gradually fills up our capture buffer as data comes in. We need to specify a *data request* signal (DREQ) so that the DMA transfers data at the proper rate.

Data request signals

The DMA can transfer data incredibly fast, and almost invariably this will be much faster than your PIO program actually needs. The DMA paces itself based on a data request handshake with the state machine, so there's no worry about it overflowing or underflowing a FIFO, as long as you have selected the correct DREQ signal. The state machine coordinates with the DMA to tell it when it has room available in its TX FIFO, or data available in its RX FIFO.

We need to provide the DMA channel with an initial read address, an initial write address, and the total number of reads/writes to be performed (*not* the total number of bytes). We start the DMA channel immediately – from this point

on, the DMA is poised, waiting for the state machine to produce data. As soon as data appears in the RX FIFO, the DMA will pounce and whisk the data away to our capture buffer in system memory.

As things stand right now, the state machine will immediately go into a 1-cycle loop of in instructions once enabled. Since the system memory available for capture is quite limited, it would be better for the state machine to wait for some trigger before it starts sampling. Specifically, we are using a wait pin instruction to stall the state machine until a certain pin goes high or low, and again we are using one of the pio_encode_ functions to encode this instruction on-the-fly.

pio_sm_exec tells the state machine to immediately execute some instruction you give it. This instruction never gets written to the instruction memory, and if the instruction stalls (as it will in this case — a wait instruction's job is to stall) then the state machine will latch the instruction until it completes. With the state machine stalled on the wait instruction, we can enable it without being immediately flooded by data.

At this point everything is armed and waiting for the trigger signal from the chosen GPIO. This will lead to the following sequence of events:

- 1. The wait instruction will clear
- 2. On the very next cycle, state machine will start to execute in instructions from the program memory
- 3. As soon as data appears in the RX FIFO, the DMA will start to transfer it.
- 4. Once the requested amount of data has been transferred by the DMA, it'll automatically stop

State Machine EXEC Functionality

So far our state machines have executed instructions from the instruction memory, but there are other options. One is the SMx_INSTR register (used by pio_sm_exec()): the state machine will immediately execute whatever you write here, momentarily interrupting the current program it's running if necessary. This is useful for poking around inside the state machine from the system side, for initial setup.

The other two options, which use the same underlying hardware, are out exec (shift out an instruction from the data being streamed through the OSR, and execute it) and mov exec (execute an instruction stashed in e.g. a scratch register). Besides making people's eyes bulge, these are really useful if you want the state machine to perform some data-defined operation at a certain point in an output stream.

The example code provides this cute function for displaying the captured logic trace as ASCII art in a terminal:

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/logic_analyser/logic_analyser.c Lines 89 - 108

```
89 void print_capture_buf(const uint32_t *buf, uint pin_base, uint pin_count, uint32_t
   n_samples) {
90
     // Display the capture buffer in text form, like this:
       // 00: __--__-
 92
       // 01: ____-
       printf("Capture:\n");
93
       // Each FIFO record may be only partially filled with bits, depending on
94
95
       // whether pin_count is a factor of 32.
       uint record_size_bits = bits_packed_per_word(pin_count);
96
       for (int pin = 0; pin < pin_count; ++pin) {</pre>
97
           printf("%02d: ", pin + pin_base);
98
           for (int sample = 0; sample < n_samples; ++sample) {</pre>
99
100
               uint bit_index = pin + sample * pin_count;
101
               uint word_index = bit_index / record_size_bits;
               // Data is left-justified in each FIFO entry, hence the (32 - record_size_bits)
   offset
103
               uint word_mask = 1u << (bit_index % record_size_bits + 32 - record_size_bits);</pre>
104
                printf(buf[word_index] & word_mask ? "-" : "_");
105
           }
106
           printf("\n");
107
       }
108 }
```

We have everything we need now for RP2040 to capture a logic trace of its own pins, whilst running some other program. Here we're setting up a PWM slice to output at around 15MHz on two GPIOs, and attaching our brand spanking new logic analyser to those same two GPIOs.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/logic_analyser/logic_analyser.c Lines 110 - 159

```
110 int main() {
111
       stdio_init_all();
112
        printf("PIO logic analyser example\n");
113
114
       // We're going to capture into a u32 buffer, for best DMA efficiency. Need
115
       // to be careful of rounding in case the number of pins being sampled
116
        // isn't a power of 2.
        uint total_sample_bits = CAPTURE_N_SAMPLES * CAPTURE_PIN_COUNT;
117
        total_sample_bits += bits_packed_per_word(CAPTURE_PIN_COUNT) - 1;
118
        uint buf_size_words = total_sample_bits / bits_packed_per_word(CAPTURE_PIN_COUNT);
119
        uint32_t *capture_buf = malloc(buf_size_words * sizeof(uint32_t));
120
121
        hard_assert(capture_buf);
122
123
        // Grant high bus priority to the DMA, so it can shove the processors out
124
        // of the way. This should only be needed if you are pushing things up to
125
        // >16bits/clk here, i.e. if you need to saturate the bus completely.
        bus_ctrl_hw->priority = BUSCTRL_BUS_PRIORITY_DMA_W_BITS |
126
    BUSCTRL_BUS_PRIORITY_DMA_R_BITS;
127
        PIO pio = pio0;
128
        uint sm = 0;
129
130
        uint dma_chan = 0;
131
132
        logic_analyser_init(pio, sm, CAPTURE_PIN_BASE, CAPTURE_PIN_COUNT, 1.f);
133
134
        printf("Arming trigger\n");
        logic_analyser_arm(pio, sm, dma_chan, capture_buf, buf_size_words, CAPTURE_PIN_BASE,
135
    true);
136
        printf("Starting PWM example\n");
137
138
        // PWM example: -----
        gpio_set_function(CAPTURE_PIN_BASE, GPIO_FUNC_PWM);
139
140
        gpio_set_function(CAPTURE_PIN_BASE + 1, GPIO_FUNC_PWM);
141
        // Topmost value of 3: count from 0 to 3 and then wrap, so period is 4 cycles
142
        pwm_hw->slice[0].top = 3;
        // Divide frequency by two to slow things down a little
143
144
        pwm_hw->slice[0].div = 4 << PWM_CH0_DIV_INT_LSB;</pre>
145
        // Set channel A to be high for 1 cycle each period (duty cycle 1/4) and
146
        // channel B for 3 cycles (duty cycle 3/4)
147
        pwm_hw->slice[0].cc =
148
                (1 << PWM_CH0_CC_A_LSB) |
                (3 << PWM_CH0_CC_B_LSB);</pre>
149
        // Enable this PWM slice
150
151
        pwm_hw->slice[0].csr = PWM_CH0_CSR_EN_BITS;
152
153
        // The logic analyser should have started capturing as soon as it saw the
154
155
        // first transition. Wait until the last sample comes in from the DMA.
156
        dma_channel_wait_for_finish_blocking(dma_chan);
157
        print_capture_buf(capture_buf, CAPTURE_PIN_BASE, CAPTURE_PIN_COUNT, CAPTURE_N_SAMPLES);
158
159 }
```

The output of the program looks like this:

3.2.4. Further examples

Hopefully what you have seen so far has given some idea of how PIO applications can be built with the SDK. The RP2040 Datasheet contains *many* more documented examples, which highlight particular hardware features of PIO, or show how particular hardware interfaces can be implemented.

You can also browse the pio/ directory in the Pico Examples repository.

3.3. Using PIOASM, the PIO Assembler

Up until now, we have glossed over the details of how the assembly program in our .pio file is translated into a binary program, ready to be loaded into our PIO state machine. Programs that handle this task — translating assembly code into binary — are generally referred to as assemblers, and PIO is no exception in this regard. The SDK includes an assembler for PIO, called pioasm. The SDK handles the details of building this tool for you behind the scenes, and then using it to build your PIO programs, for you to #include from your C or C++ program. pioasm can also be used directly, and has a few features not used by the C++ SDK, such as generating programs suitable for use with the MicroPython PIO library.

If you have built the pico-examples repository at any point, you will likely already have a pioasm binary in your build directory, located under build/tools/pioasm/pioasm, which was bootstrapped for you before building any applications that depend on it. If we want a standalone copy of pioasm, perhaps just to explore the available command-line options, we can obtain it as follows (assuming the SDK is extracted at \$PICO_SDK_PATH):

```
$ mkdir pioasm_build
$ cd pioasm_build
$ cmake $PICO_SDK_PATH/tools/pioasm
$ make
```

And then invoke as:

```
$ ./pioasm
```

3.3.1. Usage

A description of the command line arguments can be obtained by running:

```
$ pioasm -?
```

giving:

```
usage: pioasm <options> <input> (<output>)
Assemble file of PIO program(s) for use in applications.
<input>
                   the input filename
<output>
                   the output filename (or filename prefix if the output
                        format produces multiple outputs).
                    if not specified, a single output will be written to stdout
options:
-o <output_format>
                   select output_format (default 'c-sdk'); available options are:
                            C header suitable for use with the Raspberry Pi Pico SDK
                        python
                            Python file suitable for use with MicroPython
                        hex
                            Raw hex output (only valid for single program inputs)
                    add a parameter to be passed to the outputter
-p <output_param>
-?, --help
                     print this help and exit
```

NOTE

Within the SDK you do not need to invoke pioasm directly, as the CMake function pico_generate_pio_header(TARGET PIO_FILE) takes care of invoking pioasm and adding the generated header to the include path of the target TARGET for you.

3.3.2. Directives

The following directives control the assembly of PIO programs:

Table 5. pioasm directives

.define (PUBLIC) <symbol> <value>

Define an integer symbol named <symbol> with the value <value> (see Section 3.3.3). If this .define appears before the first program in the input file, then the define is global to all programs, otherwise it is local to the program in which it occurs. If PUBLIC is specified the symbol will be emitted into the assembled output for use by user code. For the SDK this takes the form of:

#define csymbol> value for program symbols or #define <symbol> value for global symbols

Start a new program with the name < name >. Note that that name is used in .program <name>

code so should be alphanumeric/underscore not starting with a digit. The program lasts until another .program directive or the end of the source file. PIO

instructions are only allowed within a program

.origin <offset> Optional directive to specify the PIO instruction memory offset at which the program must load. Most commonly this is used for programs that must load

at offset 0, because they use data based JMPs with the (absolute) jmp target being stored in only a few bits. This directive is invalid outside of a program

.side_set <count> (opt) (pindirs) If this directive is present, <count> indicates the number of side-set bits to be

> used. Additionally opt may be specified to indicate that a side <value> is optional for instructions (note this requires stealing an extra bit — in addition to the <count> bits - from those available for the instruction delay). Finally, pindirs may be specified to indicate that the side set values should be applied to the PINDIRs and not the PINs. This directive is only valid within a program

before the first instruction

Place prior to an instruction, this directive specifies the instruction where .wrap_target execution continues due to program wrapping. This directive is invalid outside of a program, may only be used once within a program, and if not specified defaults to the start of the program .wrap Placed after an instruction, this directive specifies the instruction after which, in normal control flow (i.e. jmp with false condition, or no jmp), the program wraps (to .wrap_target instruction). This directive is invalid outside of a program, may only be used once within a program, and if not specified defaults to after the last program instruction. .lang_opt <lang> <name> <option> Specifies an option for the program related to a particular language generator. (See Section 3.3.10). This directive is invalid outside of a program .word <value> Stores a raw 16-bit value as an instruction in the program. This directive is invalid outside of a program.

3.3.3. Values

The following types of values can be used to define integer numbers or branch targets

Table 6. Values in pioasm. i.e. <value>

integer	An integer value e.g. 3 or -7
hex	A hexadecimal value e.g. 0xf
binary	A binary value e.g. 0b1001
symbol	A value defined by a .define (see [pioasm_define])
<label></label>	The instruction offset of the label within the program. This makes most sense when used with a JMP instruction (see Section 3.4.2)
(<expression>)</expression>	An expression to be evaluated; see expressions. Note that the parentheses are necessary.

3.3.4. Expressions

Expressions may be freely used within pioasm values.

Table 7. Expressions in pioasm i.e. <expression>

<expression> + <expression></expression></expression>	The sum of two expressions
<expression> - <expression></expression></expression>	The difference of two expressions
<expression> * <expression></expression></expression>	The multiplication of two expressions
<expression> / <expression></expression></expression>	The integer division of two expressions
- <expression></expression>	The negation of another expression
:: <expression></expression>	The bit reverse of another expression
<value></value>	Any value (see Section 3.3.3)

3.3.5. Comments

Line comments are supported with // or ;

C-style block comments are supported via /* and */

3.3.6. Labels

Labels are of the form:

<svmbol>:

or

PUBLIC <symbol>:

at the start of a line.



TIP

A label is really just an automatic .define with a value set to the current program instruction offset. A PUBLIC label is exposed to the user code in the same way as a PUBLIC .define.

3.3.7. Instructions

All pioasm instructions follow a common pattern:

<instruction> (side <side_set_value>) ([<delay_value>])

where:

<instruction>

Is an assembly instruction detailed in the following sections. (See Section 3.4)

<side_set_value>

Is a value (see Section 3.3.3) to apply to the side_set pins at the start of the instruction. Note that the rules for a side-set value via side <side_set_value> are dependent on the .side_set (see [pioasm_side_set]) directive for the program. If no .side_set is specified then the side <side_set_value> is invalid, if an optional number of sideset pins is specified then side <side_set_value> may be present, and if a non-optional number of sideset pins is specified, then side <side_set_value> is required. The <side_set_value> must fit within the number of side-set bits specified in the .side_set directive.

<delay_value>

Specifies the number of cycles to delay after the instruction completes. The delay_value is specified as a value (see Section 3.3.3), and in general is between 0 and 31 inclusive (a 5-bit value), however the number of bits is reduced when sideset is enabled via the .side_set (see [pioasm_side_set]) directive. If the <delay_value> is not present, then the instruction has no delay



NOTE

pioasm instruction names, keywords and directives are case insensitive; lower case is used in the Assembly Syntax sections below as this is the style used in the SDK.



NOTE

Commas appear in some Assembly Syntax sections below, but are entirely optional, e.g. out pins, 3 may be written out pins 3, and jmp x-- label may be written as jmp x--, label. The Assembly Syntax sections below uses the first style in each case as this is the style used in the SDK.

3.3.8. Pseudoinstructions

Currently pioasm provides one pseudoinstruction, as a convenience:

nop Assembles to mov y, y. "No operation", has no particular side effect, but a useful vehicle for a side-set operation or an extra delay.

3.3.9. Output pass through

Text in the PIO file may be passed, unmodified, to the output based on the language generator being used.

For example the following (comment and function) would be included in the generated header when the default c-sdk language generator is used.

```
% c-sdk {

// an inline function (since this is going in a header file)
static inline int some_c_code() {
    return 0;
}
```

The general format is

```
% target {
pass through contents
%}
```

with targets being recognized by a particular language generator (see Section 3.3.10; note that target is usually the language generator name e.g. c-sdk, but could potentially be some_language.some_group if the language generator supports different classes of pass through with different output locations.

This facility allows you to encapsulate both the PIO program and the associated setup required in the same source file. See Section 3.3.10 for a more complete example.

3.3.10. Language generators

The following example shows a multi program source file (with multiple programs) which we will use to highlight c-sdk and python output features

 ${\it Pico\ Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/ws2812.pio.pico-examples.pico-examples.pico-e$

```
1;
2 ; Copyright (c) 2020 Raspberry Pi (Trading) Ltd.
3;
4 ; SPDX-License-Identifier: BSD-3-Clause
5;
6
7 .program ws2812
8 .side_set 1
10 .define public T1 2
11 .define public T2 5
12 .define public T3 3
13
14 .lang_opt python sideset_init = pico.PIO.OUT_HIGH
15 .lang_opt python out_init = pico.PIO.OUT_HIGH
16 .lang_opt python out_shiftdir = 1
17
```

```
18 .wrap_target
19 bitloop:
20 out x, 1
                   side 0 [T3 - 1] ; Side-set still takes place when instruction stalls
21
      jmp !x do_zero side 1 [T1 - 1] ; Branch on the bit we shifted out. Positive pulse
22 do one:
23 jmp bitloop side 1 [T2 - 1]; Continue driving high, for a long pulse
24 do_zero:
25 nop
                    side 0 [T2 - 1] ; Or drive low, for a short pulse
26 .wrap
27
28 % c-sdk {
29 #include "hardware/clocks.h"
31 static inline void ws2812_program_init(PIO pio, uint sm, uint offset, uint pin, float freq,
  bool rgbw) {
32
33
    pio_gpio_init(pio, pin);
34
    pio_sm_set_consecutive_pindirs(pio, sm, pin, 1, true);
   pio_sm_config c = ws2812_program_get_default_config(offset);
37
   sm_config_set_sideset_pins(&c, pin);
38
    sm_config_set_out_shift(&c, false, true, rgbw ? 32 : 24);
     sm_config_set_fifo_join(&c, PIO_FIFO_JOIN_TX);
39
40
      int cycles_per_bit = ws2812_T1 + ws2812_T2 + ws2812_T3;
41
      float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit);
42
43
      sm_config_set_clkdiv(&c, div);
44
45
      pio_sm_init(pio, sm, offset, &c);
46
      pio_sm_set_enabled(pio, sm, true);
47 }
48 %}
49
50 .program ws2812_parallel
51
52 .define public T1 2
53 .define public T2 5
54 .define public T3 3
56 .wrap_target
57 out x, 32
mov pins, !null [T1-1]
59 mov pins, x [T2-1]
60 mov pins, null [T3-2]
61 .wrap
62
63 % c-sdk {
64 #include "hardware/clocks.h"
65
66 static inline void ws2812_parallel_program_init(PIO pio, uint sm, uint offset, uint
  pin_base, uint pin_count, float freq) {
67
       for(uint i=pin_base; i<pin_base+pin_count; i++) {</pre>
68
          pio_gpio_init(pio, i);
69
70
       pio_sm_set_consecutive_pindirs(pio, sm, pin_base, pin_count, true);
71
72
      pio_sm_config c = ws2812_parallel_program_get_default_config(offset);
73
      sm_config_set_out_shift(&c, true, true, 32);
74
      sm_config_set_out_pins(&c, pin_base, pin_count);
75
     sm_config_set_set_pins(&c, pin_base, pin_count);
76
     sm_config_set_fifo_join(&c, PIO_FIFO_JOIN_TX);
77
78
      int cycles_per_bit = ws2812_parallel_T1 + ws2812_parallel_T2 + ws2812_parallel_T3;
```

```
79
      float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit);
80
   sm_config_set_clkdiv(&c, div);
81
82 pio_sm_init(pio, sm, offset, &c);
    pio_sm_set_enabled(pio, sm, true);
83
84 }
85 %}
```

3.3.10.1. c-sdk

The c-sdk language generator produces a single header file with all the programs in the PIO source file:

The pass through sections (% c-sdk {) are embedded in the output, and the PUBLIC defines are available via #define



pioasm creates a function for each program (e.g. ws2812_program_get_default_config()) returning a pio_sm_config based on the .side_set, .wrap and .wrap_target settings of the program, which you can then use as a basis for configuration the PIO state machine.

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/generated/ws2812.pio.h

```
1 // ----- //
2 // This file is autogenerated by pioasm; do not edit! //
4
5 #pragma once
7 #if !PICO_NO_HARDWARE
8 #include "hardware/pio.h"
9 #endif
10
11 // ----- //
12 // ws2812 //
13 // ----- //
14
15 #define ws2812_wrap_target 0
16 #define ws2812_wrap 3
17
18 #define ws2812_T1 2
19 #define ws2812_T2 5
20 #define ws2812_T3 3
21
22 static const uint16_t ws2812_program_instructions[] = {
23
     // .wrap_target
                                    side 0 [2]
24
   0x6221, // 0: out x, 1
                                    side 1 [1]
25
   0x1123, // 1: jmp !x, 3
   0x1400, // 2: jmp 0
                                    side 1 [4]
26
27
    0xa442, // 3: nop
                                     side 0 [4]
         // .wrap
28
29 };
31 #if !PICO_NO_HARDWARE
32 static const struct pio_program ws2812_program = {
.instructions = ws2812_program_instructions,
34
      .length = 4,
35
     .origin = -1,
36 };
37
```

```
38 static inline pio_sm_config ws2812_program_get_default_config(uint offset) {
39 pio_sm_config c = pio_get_default_sm_config();
40
      sm_config_set_wrap(&c, offset + ws2812_wrap_target, offset + ws2812_wrap);
41
      sm_config_set_sideset(&c, 1, false, false);
42
      return c:
43 }
44
45 #include "hardware/clocks.h"
46 static inline void ws2812_program_init(PIO pio, uint sm, uint offset, uint pin, float freq,
  bool rgbw) {
47
     pio_gpio_init(pio, pin);
48
      pio_sm_set_consecutive_pindirs(pio, sm, pin, 1, true);
      pio_sm_config c = ws2812_program_get_default_config(offset);
49
50
      sm_config_set_sideset_pins(&c, pin);
51
      sm_config_set_out_shift(&c, false, true, rgbw ? 32 : 24);
52
      sm_config_set_fifo_join(&c, PIO_FIFO_JOIN_TX);
53
      int cycles_per_bit = ws2812_T1 + ws2812_T2 + ws2812_T3;
      float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit);
55
      sm_config_set_clkdiv(&c, div);
      pio_sm_init(pio, sm, offset, &c);
57
      pio_sm_set_enabled(pio, sm, true);
58 }
59
60 #endif
61
62 // ----- //
63 // ws2812_parallel //
64 // ----- //
66 #define ws2812_parallel_wrap_target 0
67 #define ws2812_parallel_wrap 3
69 #define ws2812_parallel_T1 2
70 #define ws2812_parallel_T2 5
71 #define ws2812_parallel_T3 3
73 static const uint16_t ws2812_parallel_program_instructions[] = {
74
          // .wrap_target
    0x6020, // 0: out x, 32
75
76 0xa10b, // 1: mov pins, !null
                                               [1]
77
    0xa401, // 2: mov pins, x
                                                [4]
78
   0xa103, // 3: mov pins, null
                                                [1]
79
            // .wrap
80 };
81
82 #if !PICO_NO_HARDWARE
83 static const struct pio_program ws2812_parallel_program = {
84
    .instructions = ws2812_parallel_program_instructions,
85
      .lenath = 4.
86
      .origin = -1,
87 };
89 static inline pio_sm_config ws2812_parallel_program_get_default_config(uint offset) {
90
      pio_sm_config c = pio_get_default_sm_config();
91
      sm_config_set_wrap(&c, offset + ws2812_parallel_wrap_target, offset +
  ws2812_parallel_wrap);
92
    return c;
93 }
94
95 #include "hardware/clocks.h"
96 static inline void ws2812_parallel_program_init(PIO pio, uint sm, uint offset, uint
  pin_base, uint pin_count, float freq) {
97 for(uint i=pin_base; i<pin_base+pin_count; i++) {
```

```
98
            pio_gpio_init(pio, i);
99
       }
100
       pio_sm_set_consecutive_pindirs(pio, sm, pin_base, pin_count, true);
       pio_sm_config c = ws2812_parallel_program_get_default_config(offset);
101
       sm_config_set_out_shift(&c, true, true, 32);
102
       sm_config_set_out_pins(&c, pin_base, pin_count);
103
104
       sm_config_set_set_pins(&c, pin_base, pin_count);
105
       sm_config_set_fifo_join(&c, PIO_FIFO_JOIN_TX);
106
       int cycles_per_bit = ws2812_parallel_T1 + ws2812_parallel_T2 + ws2812_parallel_T3;
107
       float div = clock_get_hz(clk_sys) / (freq * cycles_per_bit);
108
       sm_config_set_clkdiv(&c, div);
109
       pio_sm_init(pio, sm, offset, &c);
110
       pio_sm_set_enabled(pio, sm, true);
111 }
112
113 #endif
```

3.3.10.2. python

The python language generator produces a single python file with all the programs in the PIO source file:

The pass through sections (% python {) would be embedded in the output, and the PUBLIC defines are available as python

Also note the use of .lang_opt python to pass initializers for the @pico.asm_pio decorator



The python language output is provided as a utility. MicroPython supports programming with the PIO natively, so you may only want to use pioasm when sharing PIO code between the SDK and MicroPython. No effort is currently made to preserve label names, symbols or comments, as it is assumed you are either using the PIO file as a source or python; not both. The python language output can of course be used to bootstrap your MicroPython PIO development based on an existing PIO file.

 ${\it Pico~Examples:} \ https://github.com/raspberrypi/pico-examples/blob/master/pio/ws2812/generated/ws2812.py$

```
2 # This file is autogenerated by pioasm; do not edit! #
3 # -----
4
5 import rp2
6 from machine import Pin
 7 # ----- #
8 # ws2812 #
9 # ----- #
10
11 \text{ ws} 2812\_T1 = 2
12 \text{ ws} 2812\_T2 = 5
13 \text{ ws} 2812 \text{ T3} = 3
14
15 @rp2.asm_pio(sideset_init=pico.PIO.OUT_HIGH, out_init=pico.PIO.OUT_HIGH, out_shiftdir=1)
16 def ws2812():
17
    wrap_target()
18
      label("0")
                              .side(\theta) [2] # \theta
19
      out(x, 1)
                              .side(1) [1] # 1
20
      jmp(not_x, "3")
21
      jmp("0")
                               .side(1) [4] # 2
22
      label("3")
                               .side(0) [4] # 3
23
      nop()
```

```
24
     wrap()
25
26
27
28 # ----- #
29 # ws2812_parallel #
32 \text{ ws} 2812\_parallel\_T1 = 2
33 ws2812_parallel_T2 = 5
34 \text{ ws} 2812\_parallel\_T3 = 3
35
36 @rp2.asm_pio()
37 def ws2812_parallel():
38 wrap_target()
39 out(x, 32)
                                                # 0
39 out(x, 32) # 0
40 mov(pins, invert(null)) [1] # 1
41 mov(pins, x) [4] # 2
42 mov(pins, null) [1] # 3
43 wrap()
```

3.3.10.3. hex

The hex generator only supports a single input program, as it just dumps the raw instructions (one per line) as a 4-character hexadecimal number.

Given:

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/squarewave/squarewave.pio

```
1 ;
2 ; Copyright (c) 2020 Raspberry Pi (Trading) Ltd.
3 ;
4 ; SPDX-License-Identifier: BSD-3-Clause
5 ;
6
7 .program squarewave
8 set pindirs, 1 ; Set pin to output
9 again:
10 set pins, 1 [1] ; Drive pin high and then delay for one cycle
11 set pins, 0 ; Drive pin low
12 jmp again ; Set PC to label `again`
```

The hex output produces:

Pico Examples: https://github.com/raspberrypi/pico-examples/blob/master/pio/squarewave/generated/squarewave.hex

```
1 e081
2 e101
3 e000
4 0001
```

3.4. PIO Instruction Set Reference

NOTE

This section refers in places to concepts and pieces of hardware discussed in the RP2040 Datasheet. You are encouraged to read the PIO chapter of the datasheet to get the full context for what these instructions do.

3.4.1. **Summary**

PIO instructions are 16 bits long, and have the following encoding:

Table 8. PIO instruction encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
JMP	0	0	0		Del	ay/side	-set		C	Conditio	n			Addres	S	
WAIT	0	0	1		Del	ay/side	-set		Pol	Sou	ırce			Index		
IN	0	1	0		Del	ay/side	-set			Source)		E	Bit cour	nt	
OUT	0	1	1		Del	ay/side	-set		De	estinati	on		E	Bit cour	nt	
PUSH	1	0	0		Del	ay/side	-set		0	IfF	Blk	0	0	0	0	0
PULL	1	0	0		Del	ay/side	-set		1	IfE	Blk	0	0	0	0	0
MOV	1	0	1		Del	ay/side	-set		De	estinati	on	С)p		Source	
IRQ	1	1	0		Del	ay/side	-set		0	Clr	Wait			Index		
SET	1	1	1		Del	ay/side	-set		De	estinati	on			Data		

All PIO instructions execute in one clock cycle.

The Delay/side-set field is present in all instructions. Its exact use is configured for each state machine by PINCTRL_SIDESET_COUNT:

- Up to 5 MSBs encode a side-set operation, which optionally asserts a constant value onto some GPIOs, concurrently with main instruction execution logic
- Remaining LSBs (up to 5) encode the number of idle cycles inserted between this instruction and the next

3.4.2. JMP

3.4.2.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
JMP	0	0	0		Del	ay/side	-set		С	Conditio	n		,	Addres	S	

3.4.2.2. Operation

Set program counter to Address if Condition is true, otherwise no operation.

Delay cycles on a JMP always take effect, whether Condition is true or false, and they take place after Condition is evaluated and the program counter is updated.

- Condition:
 - o 000: (no condition): Always
 - o 001: !X: scratch X zero

o 010: X--: scratch X non-zero, prior to decrement

o 011: !Y: scratch Y zero

o 100: Y--: scratch Y non-zero, prior to decrement

o 101: X!=Y: scratch X not equal scratch Y

o 110: PIN: branch on input pin

o 111: !OSRE: output shift register not empty

 Address: Instruction address to jump to. In the instruction encoding this is an absolute address within the PIO instruction memory.

JMP PIN branches on the GPIO selected by EXECCTRL_JMP_PIN, a configuration field which selects one out of the maximum of 32 GPIO inputs visible to a state machine, independently of the state machine's other input mapping. The branch is taken if the GPIO is high.

!OSRE compares the bits shifted out since the last PULL with the shift count threshold configured by SHIFTCTRL_PULL_THRESH. This is the same threshold used by autopull.

JMP X-- and JMP Y-- always decrement scratch register X or Y, respectively. The decrement is not conditional on the current value of the scratch register. The branch is conditioned on the *initial* value of the register, i.e. before the decrement took place: if the register is initially nonzero, the branch is taken.

3.4.2.3. Assembler Syntax

jmp (<cond>) <target>

where:

<cond>

Is an optional condition listed above (e.g. !x for scratch X zero). If a condition code is not specified, the branch is always taken

<target>

Is a program label or value (see Section 3.3.3) representing instruction offset within the program (the first instruction being offset 0). Note that because the PIO JMP instruction uses absolute addresses in the PIO instruction memory, JMPs need to be adjusted based on the program load offset at runtime. This is handled for you when loading a program with the SDK, but care should be taken when encoding JMP instructions for use by OUT EXEC

3.4.3. WAIT

3.4.3.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
WAIT	0	0	1		Del	ay/side	-set		Pol	Sou	ırce			Index		

3.4.3.2. Operation

Stall until some condition is met.

Like all stalling instructions, delay cycles begin after the instruction *completes*. That is, if any delay cycles are present, they do not begin counting until *after* the wait condition is met.

Polarity:

- o 1: wait for a 1.
- o 0: wait for a 0.
- · Source: what to wait on. Values are:
 - 00: 6PIO: System GPIO input selected by Index. This is an absolute GPIO index, and is not affected by the state machine's input IO mapping.
 - 01: PIN: Input pin selected by Index. This state machine's input IO mapping is applied first, and then Index selects which of the mapped bits to wait on. In other words, the pin is selected by adding Index to the PINCTRL_IN_BASE configuration, modulo 32.
 - o 10: IRQ: PIO IRQ flag selected by Index
 - o 11: Reserved
- · Index: which pin or bit to check.

WAIT x IRQ behaves slightly differently from other WAIT sources:

- If Polarity is 1, the selected IRQ flag is cleared by the state machine upon the wait condition being met.
- The flag index is decoded in the same way as the IRQ index field: if the MSB is set, the state machine ID (0...3) is added to the IRQ index, by way of modulo-4 addition on the two LSBs. For example, state machine 2 with a flag value of '0x11' will wait on flag 3, and a flag value of '0x13' will wait on flag 1. This allows multiple state machines running the same program to synchronise with each other.

A CAUTION

WAIT 1 IRQ x should not be used with IRQ flags presented to the interrupt controller, to avoid a race condition with a system interrupt handler

3.4.3.3. Assembler Syntax

```
wait <polarity> gpio <gpio_num>
wait <polarity> pin <pin_num>
wait <polarity> irq <irq_num> ( rel )
where:
```

<polarity> Is a value (see Section 3.3.3) specifying the polarity (either 0 or 1)

Is a value (see Section 3.3.3) specifying the input pin number (as mapped by the SM input pin

mapping)

<gpio_num> Is a value (see Section 3.3.3) specifying the actual GPIO pin number

<irq_num> (rel)
Is a value (see Section 3.3.3) specifying The irq number to wait on (0-7). If rel is present, then the

actual irq number used is calculating by replacing the low two bits of the irq number (irq_num_{10}) with the low two bits of the sum ($irq_num_{10} + sm_num_{10}$) where sm_num_{10} is the state machine

number

3.4.4. IN

3.4.4.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
IN	0	1	0		Dela	ay/side	-set			Source	!		E	Bit cour	it	

3.4.4.2. Operation

Shift Bit count bits from Source into the Input Shift Register (ISR). Shift direction is configured for each state machine by SHIFTCTRL_IN_SHIFTDIR. Additionally, increase the input shift count by Bit count, saturating at 32.

- · Source:
 - o 000: PINS
 - o 001: X (scratch register X)
 - o 010: Y (scratch register Y)
 - o 011: NULL (all zeroes)
 - o 100: Reserved
 - o 101: Reserved
 - o 110: ISR
 - o 111: 0SR
- Bit count: How many bits to shift into the ISR. 1...32 bits, 32 is encoded as 00000.

If automatic push is enabled, IN will also push the ISR contents to the RX FIFO if the push threshold is reached (SHIFTCTRL_PUSH_THRESH). IN still executes in one cycle, whether an automatic push takes place or not. The state machine will stall if the RX FIFO is full when an automatic push occurs. An automatic push clears the ISR contents to all-zeroes, and clears the input shift count.

IN always uses the least significant Bit count bits of the source data. For example, if PINCTRL_IN_BASE is set to 5, the instruction IN PINS, 3 will take the values of pins 5, 6 and 7, and shift these into the ISR. First the ISR is shifted to the left or right to make room for the new input data, then the input data is copied into the gap this leaves. The bit order of the input data is not dependent on the shift direction.

NULL can be used for shifting the ISR's contents. For example, UARTs receive the LSB first, so must shift to the right. After 8 IN PINS, 1 instructions, the input serial data will occupy bits 31...24 of the ISR. An IN NULL, 24 instruction will shift in 24 zero bits, aligning the input data at ISR bits 7...0. Alternatively, the processor or DMA could perform a byte read from FIFO address + 3, which would take bits 31...24 of the FIFO contents.

3.4.4.3. Assembler Syntax

in <source>, <bit_count>

where:

<source> Is one of the sources specified above.

<bit_count>
Is a value (see Section 3.3.3) specifying the number of bits to shift (valid range 1-32)

3.4.5. OUT

3.4.5.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
OUT	0	1	1		Del	ay/side	-set		De	estinati	on		E	Bit cour	nt	

3.4.5.2. Operation

Shift Bit count bits out of the Output Shift Register (OSR), and write those bits to Destination. Additionally, increase the output shift count by Bit count, saturating at 32.

- Destination:
 - o 000: PINS
 - o 001: X (scratch register X)
 - o 010: Y (scratch register Y)
 - o 011: NULL (discard data)
 - o 100: PINDIRS
 - o 101: PC
 - o 110: ISR (also sets ISR shift counter to Bit count)
 - o 111: EXEC (Execute OSR shift data as instruction)
- Bit count: how many bits to shift out of the OSR. 1...32 bits, 32 is encoded as 00000.

A 32-bit value is written to Destination: the lower Bit count bits come from the OSR, and the remainder are zeroes. This value is the least significant Bit count bits of the OSR if SHIFTCTRL_OUT_SHIFTDIR is to the right, otherwise it is the most significant bits.

PINS and PINDIRS use the OUT pin mapping.

If automatic pull is enabled, the OSR is automatically refilled from the TX FIFO if the pull threshold, SHIFTCTRL_PULL_THRESH, is reached. The output shift count is simultaneously cleared to 0. In this case, the OUT will stall if the TX FIFO is empty, but otherwise still executes in one cycle.

OUT EXEC allows instructions to be included inline in the FIFO datastream. The OUT itself executes on one cycle, and the instruction from the OSR is executed on the next cycle. There are no restrictions on the types of instructions which can be executed by this mechanism. Delay cycles on the initial OUT are ignored, but the executee may insert delay cycles as normal.

OUT PC behaves as an unconditional jump to an address shifted out from the OSR.

3.4.5.3. Assembler Syntax

out <destination>, <bit_count>

where:

<destination> Is one of the destinations specified above.

<bit_count> Is a value (see Section 3.3.3) specifying the number of bits to shift (valid range 1-32)

3.4.6. PUSH

3.4.6.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
PUSH	1	0	0		Del	ay/side	-set		0	IfF	Blk	0	0	0	0	0

3.4.6.2. Operation

Push the contents of the ISR into the RX FIFO, as a single 32-bit word. Clear ISR to all-zeroes.

- IfFull: If 1, do nothing unless the total input shift count has reached its threshold, SHIFTCTRL_PUSH_THRESH (the same as for autopush).
- Block: If 1, stall execution if RX FIFO is full.

PUSH IFFULL helps to make programs more compact, like autopush. It is useful in cases where the IN would stall at an inappropriate time if autopush were enabled, e.g. if the state machine is asserting some external control signal at this point.

The PIO assembler sets the Block bit by default. If the Block bit is not set, the PUSH does not stall on a full RX FIFO, instead continuing immediately to the next instruction. The FIFO state and contents are unchanged when this happens. The ISR is still cleared to all-zeroes, and the FDEBUG_RXSTALL flag is set (the same as a blocking PUSH or autopush to a full RX FIFO) to indicate data was lost.

3.4.6.3. Assembler Syntax

push (iffull)

push (iffull) block

push (iffull) noblock

where:

iffull Is equivalent to Iffull == 1 above. i.e. the default if this is not specified is Iffull == 0

block Is equivalent to Block == 1 above. This is the default if neither block nor noblock are specified

noblock Is equivalent to Block == 0 above.

3.4.7. PULL

3.4.7.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
PULL	1	0	0		Dela	ay/side	-set		1	IfE	Blk	0	0	0	0	0

3.4.7.2. Operation

Load a 32-bit word from the TX FIFO into the OSR.

- IfEmpty: If 1, do nothing unless the total output shift count has reached its threshold, SHIFTCTRL_PULL_THRESH (the same as for autopull).
- Block: If 1, stall if TX FIFO is empty. If 0, pulling from an empty FIFO copies scratch X to OSR.

Some peripherals (UART, SPI...) should halt when no data is available, and pick it up as it comes in; others (I2S) should clock continuously, and it is better to output placeholder or repeated data than to stop clocking. This can be achieved with the Block parameter.

A nonblocking PULL on an empty FIFO has the same effect as MOV OSR, X. The program can either preload scratch register X with a suitable default, or execute a MOV X, OSR after each PULL NOBLOCK, so that the last valid FIFO word will be recycled until new data is available.

PULL IFEMPTY is useful if an OUT with autopull would stall in an inappropriate location when the TX FIFO is empty. For example, a UART transmitter should not stall immediately after asserting the start bit. If Empty permits some of the same program simplifications as autopull, but the stall occurs at a controlled point in the program.



NOTE

When autopull is enabled, any PULL instruction is a no-op when the OSR is full, so that the PULL instruction behaves as a barrier. OUT NULL, 32 can be used to explicitly discard the OSR contents. See the RP2040 Datasheet for more detail on autopull.

3.4.7.3. Assembler Syntax

pull (ifempty) pull (ifempty) block pull (ifempty) noblock where:

Is equivalent to IfEmpty == 1 above. i.e. the default if this is not specified is IfEmpty == 0 ifempty

block Is equivalent to Block == 1 above. This is the default if neither block nor noblock are specified

noblock Is equivalent to Block == 0 above.

3.4.8. MOV

3.4.8.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
MOV	1	0	1		Del	ay/side	-set		D€	estinati	on	О)p		Source	

3.4.8.2. Operation

Copy data from Source to Destination.

- Destination:
 - o 000: PINS (Uses same pin mapping as OUT)
 - o 001: X (Scratch register X)
 - o 010: Y (Scratch register Y)
 - o 011: Reserved
 - o 100: EXEC (Execute data as instruction)

- o 101: PC
- o 110: ISR (Input shift counter is reset to 0 by this operation, i.e. empty)
- o 111: OSR (Output shift counter is reset to 0 by this operation, i.e. full)
- Operation:
 - o 00: None
 - o 01: Invert (bitwise complement)
 - o 10: Bit-reverse
 - o 11: Reserved
- Source:
 - o 000: PINS (Uses same pin mapping as IN)
 - 。 001: X
 - o 010: Y
 - o 011: NULL
 - o 100: Reserved
 - 101: STATUS
 - o 110: ISR
 - o 111: OSR

MOV PC causes an unconditional jump. MOV EXEC has the same behaviour as OUT EXEC (Section 3.4.5), and allows register contents to be executed as an instruction. The MOV itself executes in 1 cycle, and the instruction in Source on the next cycle. Delay cycles on MOV EXEC are ignored, but the executee may insert delay cycles as normal.

The STATUS source has a value of all-ones or all-zeroes, depending on some state machine status such as FIFO full/empty, configured by EXECCTRL_STATUS_SEL.

MOV can manipulate the transferred data in limited ways, specified by the Operation argument. Invert sets each bit in Destination to the logical NOT of the corresponding bit in Source, i.e. 1 bits become 0 bits, and vice versa. Bit reverse sets each bit *n* in Destination to bit 31 - *n* in Source, assuming the bits are numbered 0 to 31.

MOV dst, PINS reads pins using the IN pin mapping, and writes the full 32-bit value to the destination without masking. The LSB of the read value is the pin indicated by PINCTRL_IN_BASE, and each successive bit comes from a higher-numbered pin, wrapping after 31.

3.4.8.3. Assembler Syntax

mov <destination>, (op) <source>

where:

<destination> Is one of the destinations specified above.

<op> If present, is:

! or ~ for NOT (Note: this is always a bitwise NOT)

:: for bit reverse

<source> Is one of the sources specified above.

3.4.9. IRQ

3.4.9.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
IRQ	1	1	0		Del	ay/side	-set		0	Clr	Wait			Index		

3.4.9.2. Operation

Set or clear the IRQ flag selected by Index argument.

- Clear: if 1, clear the flag selected by Index, instead of raising it. If Clear is set, the Wait bit has no effect.
- · Wait: if 1, halt until the raised flag is lowered again, e.g. if a system interrupt handler has acknowledged the flag.
- Index:
 - o The 3 LSBs specify an IRQ index from 0-7. This IRQ flag will be set/cleared depending on the Clear bit.
 - If the MSB is set, the state machine ID (0...3) is added to the IRQ index, by way of modulo-4 addition on the two LSBs. For example, state machine 2 with a flag value of 0x11 will raise flag 3, and a flag value of 0x13 will raise flag 1.

IRQ flags 4-7 are visible only to the state machines; IRQ flags 0-3 can be routed out to system level interrupts, on either of the PIO's two external interrupt request lines, configured by IRQ0_INTE and IRQ1_INTE.

The modulo addition bit allows relative addressing of 'IRQ' and 'WAIT' instructions, for synchronising state machines which are running the same program. Bit 2 (the third LSB) is unaffected by this addition.

If Wait is set, Delay cycles do not begin until after the wait period elapses.

3.4.9.3. Assembler Syntax

irq <irq_num> (rel)
irq set <irq_num> (rel)
irq nowait <irq_num> (rel)
irq wait <irq_num> (rel)
irq clear <irq_num> (rel)

where:

<irq_num> (rel)
Is a value (see Section 3.3.3) specifying The irq number to wait on (0-7). If rel is present, then the

actual irq number used is calculating by replacing the low two bits of the irq number (irq_num_{10}) with the low two bits of the sum ($irq_num_{10} + sm_num_{10}$) where sm_num_{10} is the state machine

number

irq Means set the IRQ without waiting

irq set Also means set the IRQ without waiting
irq nowait Again, means set the IRQ without waiting

irq wait Means set the IRQ and wait for it to be cleared before proceeding

irq clear Means clear the IRQ

3.4.10. SET

3.4.10.1. Encoding

Bit:	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
SET	1	1	1	Delay/side-set					Destination			Data				

3.4.10.2. Operation

Write immediate value Data to Destination.

- Destination:
 - o 000: PINS
 - o 001: X (scratch register X) 5 LSBs are set to Data, all others cleared to 0.
 - $\circ~$ 010: Y (scratch register Y) 5 LSBs are set to Data, all others cleared to 0.
 - o 011: Reserved
 - o 100: PINDIRS
 - o 101: Reserved
 - o 110: Reserved
 - o 111: Reserved
- Data: 5-bit immediate value to drive to pins or register.

This can be used to assert control signals such as a clock or chip select, or to initialise loop counters. As Data is 5 bits in size, scratch registers can be SET to values from 0-31, which is sufficient for a 32-iteration loop.

The mapping of SET and OUT onto pins is configured independently. They may be mapped to distinct locations, for example if one pin is to be used as a clock signal, and another for data. They may also be overlapping ranges of pins: a UART transmitter might use SET to assert start and stop bits, and OUT instructions to shift out FIFO data to the same pins.

3.4.10.3. Assembler Syntax

set <destination>, <value>

where:

<destination> Is one of the destinations specified above.

<value> The value (see Section 3.3.3) to set (valid range 0-31)